

TAV Clock Cryptography: A Stateful Cryptographic System Based on Ephemeral Structure and Continuous Physical Entropy

Carlos Alberto Terencio de Bastos

Email: caterencio@yahoo.com.br

LinkedIn: <https://www.linkedin.com/in/carlos-alberto-t-b7a055ba>

Repository: <https://github.com/carlostbastos/tav-crypto>

November 2025

Abstract

We present TAV (Transactional Asynchronous Verification), an innovative cryptographic system that integrates continuous physical entropy collection, stateful stream cipher, message authentication, and digital signatures into a unified, self-contained architecture. Unlike conventional approaches that use physical entropy solely as a seed source for mathematical constructions--such as Jitterentropy feeding SHA3-256 for conditioning--TAV embeds physical measurements throughout its entire operation using exclusively logic operations (XOR, AND, OR, bit rotation), requiring no external cryptographic primitives.

The core innovation is the concept of '**ephemeral structure**': TAV has algebraic structure (1,464 prime numbers, transactional clocks, cryptographic keys), but this structure changes unpredictably with each operation. Four independent clocks with coprime prime periods (17, 23, 31, 47) govern state transitions across six 'prime boxes'. By the Chinese Remainder Theorem, the global state cycle reaches 563,327 transactions before any repetition in clock states alone. Combined with continuous entropy injection, the effective state space exceeds 6×10^{22} configurations.

TAV provides two digital signature schemes: **Hash-Chain (Lamport-style)** with 66-byte signatures and **Commitment-Reveal** with 72-byte signatures--both dramatically smaller than ML-DSA (2,420-4,627 bytes). The system also supports self-signed certificates for identity management. A unique property of TAV is that each instance's state serves as a **device fingerprint**, enabling fraud prevention and content authenticity verification.

Version 0.9 introduces automatic encrypted checkpointing every 10,000 transactions, a threat management system with dynamic security escalation, and dead-man switch protection. Complete implementations are provided in Python, C, JavaScript, Rust, and Arduino/ESP32, with cross-platform interoperability validated through comprehensive test suites. Performance evaluation shows complete system throughput of ~500-850 KB/s depending on security level, with isolated keystream achieving ~5 MB/s--comparable to ChaCha20 in pure Python. Shannon entropy of 7.996 bits/byte (99.95% of theoretical maximum), bit bias of only 0.15%, and autocorrelation below 0.001.

TAV deliberately avoids modular addition--the "A" in the ARX paradigm used by ChaCha20, Salsa20, BLAKE, and Skein--thereby eliminating susceptibility to rotational cryptanalysis. For constrained IoT devices, TAV offers significant advantages: zero network round-trips for key establishment (stateful), minimal memory footprint (\sim 500 bytes RAM), small key sizes (16-32 bytes vs 1,184-1,568 bytes for ML-KEM), and complete self-containment without external cryptographic dependencies.

We emphasize that TAV is experimental research software without formal security proofs or independent audits. We do not claim proven security or superiority over established algorithms. Rather, we present TAV as an alternative approach with fundamentally different trade-offs, inviting community analysis. The complete source code, test vectors, and documentation are available under AGPL-3.0 license with a commercial grace period until May 2027.

Keywords: stateful stream cipher, ephemeral structure, digital signatures, device identity, physical entropy, CPU timing jitter, logic-only cryptography, post-quantum considerations, Feistel network, prime-period clocks, IoT cryptography, fraud prevention

1. Introduction

1.1 Background and Motivation

The cryptographic landscape has evolved rapidly in response to two major challenges. First, the theoretical threat of quantum computing has prompted NIST to complete its Post-Quantum Cryptography (PQC) standardization process, releasing FIPS 203 (ML-KEM), FIPS 204 (ML-DSA), and FIPS 205 (SLH-DSA) in August 2024 [1]. These standards introduce lattice-based and hash-based constructions with proven security reductions to hard mathematical problems, but at the cost of substantially larger key sizes—ML-KEM public keys range from 800 to 1,568 bytes [2].

Second, the discovery of rotational cryptanalysis by Khovratovich and Nikoli [3] demonstrated that the modular addition operation central to the ARX (Addition, Rotation, XOR) paradigm underlying ChaCha20 [4], Salsa20 [5], BLAKE [6], and Skein [7] preserves correlations between rotated pairs of values. This weakness enabled attacks on reduced-round versions of several widely-deployed ciphers, though full-round versions remain secure.

These developments motivate exploration of alternative cryptographic approaches that:

1. **Minimize mathematical structure** that quantum algorithms might exploit
2. **Avoid modular addition** and its associated correlation vulnerabilities
3. **Integrate entropy throughout** the system rather than only at initialization

4. **Remain simple and auditable** for deployment in constrained environments

1.2 The TAV Approach

TAV (Transactional Asynchronous Verification) addresses these goals through a radically different architecture:

Physics-Pervasive Design: Unlike Jitterentropy [8], which collects CPU timing jitter but then feeds it through SHA3-256 for conditioning, TAV integrates physical measurements throughout the entire cryptographic pipeline. Every encryption, decryption, key derivation, and authentication operation incorporates fresh timing measurements processed through logic-only operations.

Logic-Only Operations: TAV uses exclusively XOR, AND, OR, and bit rotation--operations that correspond directly to hardware logic gates. Modular addition is deliberately avoided. Non-linearity is achieved through AND and OR operations with carefully chosen constants, providing confusion without the rotational correlation preservation of modular addition.

Ephemeral Structure: This is the core innovation. TAV has structure (primes, clocks, keys), but this structure changes unpredictably with each operation. When an attacker attempts to model the system, the state has already changed. The structure exists but is a moving target. This differs fundamentally from both 'no structure' (chaos) and 'static structure' (traditional cryptography).

Transactional State: Four independent clocks with coprime prime periods (17, 23, 31, 47) advance with each transaction, selecting primes from six boxes and mixing state. This creates a state machine with guaranteed non-repetition for 563,327 transactions, with effective uniqueness vastly larger due to accumulated entropy.

Integrated Authentication: The MAC-Feistel construction provides message authentication using the same logic-only primitives, eliminating dependence on HMAC-SHA256 or Poly1305.

1.3 Contributions

This paper makes the following contributions:

1. **Architectural Innovation:** We introduce the concept of 'ephemeral structure'-- algebraic elements that exist but change too rapidly to be exploited. This differs from both 'no structure' (chaos) and 'static structure' (traditional cryptography). Transactional clocks with coprime prime periods govern state evolution, with cycle length mathematically guaranteed by the Chinese Remainder Theorem.
2. **Logic-Only Primitive Set:** We demonstrate that XOR, AND, OR, and bit rotation suffice for all cryptographic operations--entropy mixing, key derivation, stream generation,

message authentication, and digital signatures--without modular addition or external hash functions.

3. **Compact Digital Signatures:** Two signature schemes producing 66-72 byte signatures--36-70x smaller than post-quantum ML-DSA (2,420-4,627 bytes)--plus certificate support for identity management.
4. **Device Identity Framework:** TAV's unique state evolution enables device fingerprinting for fraud prevention, content authenticity verification, and AI-generated content detection.
5. **Checkpoint System:** Automatic encrypted state persistence every 10,000 transactions with hardware change detection, enabling recovery and continuity.
6. **Threat Management:** Automatic threat detection with dynamic security escalation and dead-man switch protection.
7. **Comprehensive Implementation:** Complete, interoperable implementations in five languages (Python, C, JavaScript, Rust, Arduino) with 100% compatibility validated through cross-platform test vectors.
8. **Extensive Validation:** Results from 40 standardized tests covering statistical randomness (NIST SP 800-22), entropy quality (SP 800-90B alignment), avalanche properties, stress conditions, and edge cases.
9. **Fair PQC Comparison:** Detailed analysis showing TAV's advantages for stateful IoT scenarios where network round-trips dominate latency.
10. **Honest Assessment:** We explicitly document limitations, including absence of formal proofs, lack of independent audit, and appropriate use cases. We position TAV as research software inviting community analysis.

1.4 Scope and Limitations

TAV is experimental software. It has not been:

- Formally verified or proven secure
- Audited by independent cryptographers
- Analyzed for side-channel vulnerabilities
- Deployed in production environments
- Subjected to extensive cryptanalytic attack

We do not recommend TAV for protecting sensitive data. Appropriate uses include:

- Research into alternative cryptographic paradigms
- Educational demonstrations of cryptographic concepts
- Experimentation in controlled environments
- Inspiration for future formally-analyzed designs

2. The Concept of Ephemeral Structure

2.1 Structure Exists, But Is a Moving Target

A critical insight distinguishes TAV from both traditional cryptography and purely chaotic systems. TAV has algebraic structure:

- 1,464 prime numbers organized in 6 boxes
- Four clocks with prime periods (17, 23, 31, 47)
- Keys derived from master entropy pool
- Constants CONST_AND and CONST_OR for non-linearity

However, the question for an attacker is: **which structure should be analyzed?**

With each transaction:

- The 4 clocks advance at coprime periods
- Prime boxes rotate in desynchronized fashion
- Fresh physical entropy is injected into the mixer
- The pool state evolves irreversibly

When the attacker attempts to model the system, the state has already changed.

2.2 Comparison: Static vs Ephemeral Structure

System	Structure	Behavior
RSA	$n = p \times q$	Static--same n forever
ECC	Curve $E(F_p)$	Static--same curve
AES	S-box, MixColumns	Static--same tables
ChaCha20	4x4 block, constants	Static per session

TAV	Primes, clocks, pool	Ephemeral--changes each operation
-----	----------------------	-----------------------------------

In TAV, even if an attacker discovers the state at instant T, at instant T+1 the system has already incorporated: new clock ticks, possible box rotations, additional physical entropy, and mixer evolution.

2.3 Three Layers of Ephemerality

TAV implements ephemerality at three distinct time scales:

Layer 1 -- Micro (each transaction): Clock counters increment, entropy pool may receive data, derivation offset changes subtly.

Layer 2 -- Medium (prime periods): Clocks fire at cycles of 17/23/31/47, boxes 1-4 rotate, prime bytes change in derivation.

Layer 3 -- Macro (hundreds/thousands tx): Boxes 5-6 rotate (every 100/1000 tx), old pool expires (TTL 1000), checkpoint saves state (every 10,000).

Each layer serves a purpose:

- Micro ensures consecutive messages have different keys
- Medium ensures long-term patterns don't repeat
- Macro ensures state recovery and stale entropy cleanup

2.4 Why This Differs from 'No Structure'

'Absence of exploitable structure' does not mean absence of structure--it means the structure is too ephemeral to be exploited. This is the difference between:

- 'There is no pattern' (chaos) -- unpredictable, non-reproducible
- 'The pattern changes before you can use it' (TAV) -- deterministic but ephemeral

TAV maintains determinism (reproducibility with same seed and operations) while making the attack window infinitesimally small. This is a genuinely different security property from conventional approaches.

3. Related Work

3.1 Physical Entropy Sources

The use of CPU timing jitter as an entropy source is well-established in the literature and deployed systems.

Jitterentropy [8], developed by Stephan Muller, has been incorporated into the Linux kernel since version 4.2 and OpenSSL 3.4. It achieves SP800-90B compliance [9] through careful measurement of CPU execution time variations. However, Jitterentropy uses SHA3-256 for entropy conditioning and serves solely as an entropy source--it requires separate implementations of cipher (e.g., ChaCha20) and MAC (e.g., HMAC or Poly1305) for complete cryptographic functionality.

The Linux kernel's /dev/random implementation [10] combines multiple entropy sources including interrupt timing, disk events, and keyboard/mouse input, then conditions them through ChaCha20-based DRNG. This represents the standard architecture: physical entropy seeds a PRNG, which then operates independently.

TAV's Distinction: Rather than using physical entropy only to seed a mathematically-defined PRNG, TAV incorporates timing measurements at every operation. The physical source is not just the seed--it permeates the ongoing state evolution.

3.2 ARX Ciphers and Rotational Cryptanalysis

The ARX paradigm--using Addition, Rotation, and XOR--underlies many efficient modern ciphers:

- **Salsa20** [5] (Bernstein, 2005): 20-round stream cipher, basis for ChaCha
- **ChaCha20** [4] (Bernstein, 2008): Modified Salsa20 with improved diffusion, standardized in RFC 8439
- **BLAKE** [6] (Aumasson et al., 2008): SHA-3 finalist using ARX structure
- **Skein** [7] (Ferguson et al., 2008): SHA-3 finalist with Threefish block cipher

These designs achieve excellent software performance (~1,000 MB/s for ChaCha20) and resist timing attacks due to constant-time operations on most platforms.

However, Khovratovich and Nikolić's rotational cryptanalysis [3] demonstrated that modular addition preserves correlations between rotated input/output pairs more strongly than XOR alone. This enabled distinguishing attacks on reduced-round versions:

- 7 rounds of Salsa20 (of 20)
- 6 rounds of ChaCha (of 20)
- Various reduced rounds of BLAKE and Skein

While full-round versions remain secure, this demonstrates a fundamental limitation of the ARX approach.

TAV's Distinction: By replacing modular addition with AND and OR operations, TAV eliminates this vulnerability class entirely. AND and OR provide non-linearity without preserving rotational correlations.

3.3 NIST Post-Quantum Cryptography Standards

In August 2024, NIST released the first three finalized post-quantum standards [1]:

FIPS 203 - ML-KEM (Module-Lattice Key Encapsulation Mechanism, derived from CRYSTALS-Kyber):

- Primary standard for key establishment
- Security based on Module Learning with Errors (MLWE) problem
- Public key sizes: 800 bytes (ML-KEM-512), 1,184 bytes (ML-KEM-768), 1,568 bytes (ML-KEM-1024)
- Ciphertext sizes: 768, 1,088, 1,568 bytes respectively

FIPS 204 - ML-DSA (Module-Lattice Digital Signature Algorithm, derived from CRYSTALS-Dilithium):

- Primary standard for digital signatures
- Signature sizes: 2,420 bytes (ML-DSA-44), 3,293 bytes (ML-DSA-65), 4,627 bytes (ML-DSA-87)
- Public key sizes: 1,312, 1,952, 2,592 bytes respectively

FIPS 205 - SLH-DSA (Stateless Hash-Based Digital Signature Algorithm, derived from SPHINCS+):

- Backup signature standard based on hash functions rather than lattices
- Larger signatures but different mathematical foundation for diversity

Additional algorithms under consideration include FN-DSA (from FALCON) and HQC (code-based KEM) [11].

These standards provide provable security reductions to hard problems but introduce:

- Dependencies on SHA3/SHAKE hash functions
- Substantially larger key and signature sizes
- Implementation complexity unsuitable for severely constrained devices

TAV's Distinction: TAV takes a fundamentally different approach. Rather than basing security on hard mathematical problems, it relies on the physical unpredictability of timing jitter and the absence of mathematical structure for quantum algorithms to exploit. This lacks formal proof but presents no obvious attack surface for Shor's or Grover's algorithms. TAV focuses on stream cipher and MAC functionality, leaving digital signatures to established algorithms (ML-DSA, ECDSA).

3.4 Stateful Encryption

Bellare, Kohno, and Shoup [12] introduced stateful public-key encryption to improve efficiency by reusing randomness across encryptions. Their framework established security definitions for stateful schemes but focused on asymmetric cryptography with simple counter-based state.

TAV's Distinction: TAV extends statefulness to symmetric cryptography with a complex multi-clock architecture. State evolution depends on:

- Transaction counts (four independent clocks)
- Prime number indices (six boxes)
- Accumulated physical entropy
- Message-dependent mixing

This creates state that is both deterministically reproducible (given the same seed and operations) and physically unpredictable (due to continuous entropy injection).

4. TAV Architecture

4.1 Design Philosophy

TAV is built on four principles:

Principle 1 - Physical Foundation: Security derives from physical phenomena (timing jitter) rather than mathematical assumptions about computational hardness.

Principle 2 - Logic Simplicity: All operations map directly to hardware logic gates (XOR, AND, OR, bit shift/rotate), enabling implementation on any processor and facilitating security analysis.

Principle 3 - Continuous Evolution: State changes with every operation, incorporating fresh entropy. There is no concept of a "static key" except at derivation time.

Principle 4 - Self-Containment: No external cryptographic libraries or primitives are required. The entire system is implemented from first principles.

4.2 System Components

TAV comprises five integrated subsystems:

4.3 Physical Entropy Generator

4.3.1 Timing Collection

The entropy generator exploits CPU execution time variations caused by:

- Cache state (hits vs. misses)
- Branch prediction accuracy
- Memory access patterns
- Instruction pipeline state
- OS scheduler interruptions
- Background system activity

```
python
```

```
def collect_timing():

    Measure execution time for simple operation.

    t1 = perf_counter_ns() # Nanosecond precision
    _ = sum(range(10))    # Simple work
    t2 = perf_counter_ns()

    return t2 - t1        # Timing in nanoseconds
```

Multiple measurements are combined via XOR to increase entropy density:

```
python
```

```
def collect_timing_xor(n_xor=2):

    Combine multiple timing measurements.

    result = 0
    for _ in range(n_xor):
        t = collect_timing()
```

```

result ^= t

result ^= (t >> 8) # Mix byte positions

result ^= (t << 3) # Additional diffusion

return result & 0xFFFFFFFFFFFFFF

```

4.3.2 Feistel Mixer

Raw timing values are processed through a Feistel network to achieve:

- Full diffusion (every input bit affects every output bit)
- Non-linearity (via AND and OR operations)
- Balance (uniform output distribution)

Optimization: Precomputed Lookup Tables

To maximize performance, rotation operations use hardcoded lookup tables. There are 8 tables of 256 bytes each, precomputed for all possible rotations (0-7 bits):

python

```

# Lookup tables for left rotation (hardcoded, not generated at runtime)

ROT_LEFT = (
    bytes([...]), # rot=0: identity
    bytes([...]), # rot=1: ((b << 1) | (b >> 7)) & 0xFF
    # ... up to rot=7
)

```

```

def feistel_mix(data: bytes, rounds: int) -> bytes:

    Mix data through Feistel network with logic-only operations.

    state = bytearray(data)
    n = len(state)

    for round in range(rounds):
        for i in range(n):

```

```

x = state[i]

# Rotation via lookup table (optimized)

x = ROT_LEFT[(round + i) & 7][x]

# AND with constant for non-linearity

x = x & CONST_AND[(i + round * 7) & 31]

# OR with constant for balance

x = x | CONST_OR[(i + round * 11) & 31]

# XOR with neighbor for avalanche

x = x ^ state[(i + round + 1) % n]

state[i] = x

return bytes(state)

```

The lookup tables occupy only 2 KB of memory (8 256 bytes) and eliminate rotation calculations at runtime, providing ~30% throughput improvement without compromising securitythe operations are mathematically identical.

The constants CONST_AND and CONST_OR are derived from prime numbers to avoid weak patterns:

python

```

CONST_AND = bytes([
    0xB7, 0x5D, 0xA3, 0xE1, 0x97, 0x4F, 0xC5, 0x2B,
    0x8D, 0x61, 0xF3, 0x1F, 0xD9, 0x73, 0x3D, 0xAF,
    0x17, 0x89, 0xCB, 0x53, 0xE7, 0x2D, 0x9B, 0x41,
    0xBB, 0x6D, 0xF1, 0x23, 0xDD, 0x7F, 0x35, 0xA9
])

```

```

CONST_OR = bytes([
    0x11, 0x22, 0x44, 0x08, 0x10, 0x21, 0x42, 0x04,
    0x12, 0x24, 0x48, 0x09, 0x14, 0x28, 0x41, 0x02,
    0x18, 0x30, 0x60, 0x05, 0x0A, 0x15, 0x2A, 0x54,
    0x19, 0x32, 0x64, 0x06, 0x0C, 0x19, 0x33, 0x66
])

```

4.4 Transactional Clock System

Four clocks operate independently with coprime prime periods:

Clock	Period	Prime Boxes Controlled
0	17	1, 2, 3
1	23	1, 3, 4
2	31	2, 3, 4
3	47	2, 4, 5

Each clock maintains a counter that increments with every transaction (encrypt, decrypt, or tick operation). When a counter reaches its period, it resets to zero and triggers a rotation in its associated prime boxes.

State Cycle Length: By the Chinese Remainder Theorem, the clock state returns to its initial configuration after:

$$\text{LCM}(17, 23, 31, 47) = 563,327 \text{ transactions}$$

However, this considers only clock counters. The actual state includes:

- Positions within each of 6 prime boxes (21 to 500 positions each)
- Accumulated entropy (64 bytes continuously modified)
- Message-dependent state evolution

The effective state space exceeds 610 distinct configurations.

4.5 Prime Box System

Six boxes contain precomputed prime numbers of increasing magnitude:

Box	Digit Range	Quantity	Examples
1	2 digits	21	11, 13, ..., 97
2	3 digits	143	101, 103, ..., 997
3	4 digits	500	1009, 1013, ..., 4993
4	5 digits	500	10007, 10009, ..., 14759
5	7 digits	200	1000003, 1000033, ..., 1002527
6	9 digits	100	100000007, ..., 100001689
Total		1,464	

Optimization: Primes as Bytes

Primes are hardcoded in two forms:

1. **As integers:** For mathematical calculations when needed
2. **As precomputed bytes:** For key derivation (8 bytes, big-endian)

python

Primes as integers

```
PRIMES_BOX_1 = [11, 13, 17, 19, 23, 29, 31, 37, 41, 43, ...]
```

Primes as bytes (precomputed once at module load)

```
PRIMES_BYTES_BOX_1 = [
    b'\x00\x00\x00\x00\x00\x00\x00\x0b', # 11
    b'\x00\x00\x00\x00\x00\x00\x00\x0d', # 13
    ...
]
```

Benefits:

- **Fast startup:** No primality testing at runtime
- **Performance:** Eliminates `to_bytes()` in each key derivation (~50% faster)
- **Deterministic behavior:** Exact sets documented and auditable

When a clock triggers rotation, its associated boxes advance their position pointer, selecting different primes for subsequent key derivation.

4.6 Security Levels

TAV supports four security levels with different parameters:

Level	Key Size	MAC Size	Nonce Size	Boxes Active	Mixer Rounds	Verify Every
IoT	128 bits	8 bytes	12 bytes	1, 2	2	20 tx
Consumer	192 bits	12 bytes	16 bytes	1, 2, 3	3	10 tx
Enterprise	256 bits	16 bytes	16 bytes	1, 2, 3, 4	4	5 tx
Military	256 bits	24 bytes	24 bytes	1, 2, 3, 4, 5, 6	6	1 tx

4.7 Checkpoint System

TAV includes automatic encrypted state persistence:

- **Automatic checkpoint** every 10,000 transactions
- **Checkpoint encrypted** by TAV itself (self-protecting)
- **Automatic restoration** on initialization
- **Hardware change detection** (warning, not blocking)
- **Transactional TTL** for entropy pool (1,000 transactions)

The checkpoint key is derived from the original seed using a separate derivation path, ensuring checkpoint security is independent of operational keys:

python

```
def _derive_checkpoint_key(self, seed) -> bytes:  
    """Derive fixed key for checkpoint - based only on seed."""  
    seed_bytes = (seed + '_TAV_CHECKPOINT_KEY_V93').encode('utf-8')  
    key = bytearray(32)  
    for i, b in enumerate(seed_bytes):  
        key[i % 32] ^= b  
    for round in range(4):
```

```

for i in range(32):
    key[i] = ROT_LEFT[(round + i) & 7][key[i]] ^ key[(i + 1) % 32]
return bytes(key)

```

State Serialization: The checkpoint includes transaction count, boot count, master entropy, clock states, box indices, hardware profile, and nonce counter. Total serialized size is approximately 200-300 bytes depending on security level.

Hardware Profile Comparison: On restoration, TAV compares the current hardware timing profile with the saved profile. A similarity below 0.7 triggers a warning but does not block operation, allowing for legitimate hardware changes while alerting to potential issues.

4.8 MAC-Feistel Authentication

Message authentication uses a Feistel construction rather than HMAC:

python

```
def mac_feistel(message: bytes, key: bytes, rounds: int) -> bytes:
```

Compute MAC using Feistel structure.

Initialize state from key

```
state = bytearray(32)
```

```
for i, k in enumerate(key):
```

```
    state[i % 32] ^= k
```

Process message in blocks

```
for i in range(0, len(message), 32):
```

```
    block = message[i:i+32].ljust(32, b'\x00')
```

```
    for j in range(32):
```

```
        state[j] ^= block[j]
```

```
    state = bytearray(feistel_mix(bytes(state), rounds))
```

Finalization: mix in length

```
length_bytes = len(message).to_bytes(8, 'big')
```

```

for i, b in enumerate(length_bytes):
    state[i] ^= b

# Final mixing rounds

state = bytarray(feistel_mix(bytes(state), rounds))

return bytes(state[:mac_size])

```

This achieves:

- No dependence on SHA-256 or any external hash
- Consistent use of logic-only operations throughout TAV
- Avalanche effect: 50.78% bit change on single-bit message modification (measured)

4.9 Digital Signatures

TAV provides two digital signature schemes, both fully implemented:

4.9.1 Hash-Chain (Lamport-style)

A stateful signature scheme based on hash chains:

Concept:

S_0 (seed) $\rightarrow S_1 \rightarrow S_2 \rightarrow \dots \rightarrow S_n$ (public key)

Each $S_i = \text{hash}(S_{\{i-1\}})$

Key Generation:

- Generate chain of n hashes from seed
- Public key = final hash (S_n)
- Private key = seed + chain position

Signing message #i:

1. Reveal $S_{\{n-i\}}$ (the appropriate chain element)
2. MAC = $\text{hash}(\text{message} \parallel S_{\{n-i\}})$
3. Signature = (index, $S_{\{n-i\}}$, MAC)

Verification:

1. Apply hash (index+1) times to revealed value
2. Result must equal public key
3. Verify MAC

Properties:

- Signature size: **66 bytes** ($2 + 32 + 32$)
- Quantum-resistant by design (no mathematical structure)
- Limited signatures per key (chain length)
- Simple, well-studied construction

c

```
tav_result_t tav_sign_chain_keygen(tav_sign_chain_t* keys,  
                                    const uint8_t* seed, size_t seed_len);  
  
tav_result_t tav_sign_chain_sign(tav_sign_chain_t* keys,  
                                  const uint8_t* message, size_t msg_len,  
                                  uint8_t* signature, size_t* sig_len);  
  
tav_result_t tav_sign_chain_verify(const uint8_t* public_key,  
                                    const uint8_t* message, size_t msg_len,  
                                    const uint8_t* signature, size_t sig_len);
```

4.9.2 Commitment-Reveal (TAV State)

A signature scheme leveraging TAV's ephemeral state:

Concept:

- Public key = $\text{hash}(\text{master_entropy})$ = "commitment"
- Signing proves knowledge of state that generates the commitment
- Uses TAV's MAC-Feistel to bind message to state

Signing:

1. Capture current tx_count
2. Derive state_proof = $\text{hash}(\text{master_entropy} \parallel \text{tx_count})$

3. Derive sign_key = hash(state_proof)
4. MAC = hash(message || tx_count) XOR sign_key
5. Signature = (tx_count, state_proof, MAC)

Properties:

- Signature size: **72 bytes** (8 + 32 + 32)
- Unlimited signatures (uses dynamic state)
- Signatures are unique per transaction
- Requires state synchronization for verification

C

```
tav_result_t tav_sign_commit_keygen(tav_sign_commit_t* keys,
                                     const uint8_t* seed, size_t seed_len,
                                     tav_level_t level);

tav_result_t tav_sign_commit_sign(tav_sign_commit_t* keys,
                                   const uint8_t* message, size_t msg_len,
                                   uint8_t* signature, size_t* sig_len);

tav_result_t tav_sign_commit_verify(const uint8_t* public_commitment,
                                    const uint8_t* message, size_t msg_len,
                                    const uint8_t* signature, size_t sig_len);
```

4.9.3 Certificates

TAV supports X.509-style certificates for identity binding:

C

```
typedef struct {

    char identity[TAV_CERT_MAX_IDENTITY];
    uint8_t public_key[TAV_SIGN_HASH_SIZE];
    uint64_t valid_from;
    uint64_t valid_until;
```

```

        uint8_t issuer_sig[TAV_SIGN_MAX_SIG];
        uint8_t issuer_sig_len;
    } tav_cert_t;

tav_result_t tav_cert_create_self_signed(tav_cert_t* cert,
                                         const char* identity,
                                         tav_sign_chain_t* keys,
                                         uint64_t validity_seconds);

tav_result_t tav_cert_verify(const tav_cert_t* cert,
                            const uint8_t* issuer_public_key);

```

4.9.4 Signature Size Comparison

Algorithm	Signature Size	Public Key Size	Type
TAV Hash-Chain	66 bytes	32 bytes	Stateful
TAV Commitment	72 bytes	32 bytes	Stateful
ECDSA P-256	64 bytes	64 bytes	Stateless
Ed25519	64 bytes	32 bytes	Stateless
ML-DSA-44	2,420 bytes	1,312 bytes	Stateless
ML-DSA-65	3,293 bytes	1,952 bytes	Stateless
ML-DSA-87	4,627 bytes	2,592 bytes	Stateless

TAV signatures are **36-70x smaller** than post-quantum ML-DSA signatures.

4.10 Threat Management System

TAV includes an automatic threat detection and response system:

4.10.1 Threat Detection

The system monitors for:

- **Consecutive MAC failures:** Potential tampering or replay attack
- **Hardware profile changes:** Possible device cloning or migration

- **Abnormal transaction rates:** Potential brute-force attempts

4.10.2 Dynamic Security Escalation

When threats are detected, TAV automatically escalates security level:

python

```
class GerenciadorAmeacas:

    def registrar_ameaca(self, tipo: str, severidade: int, detalhes: str):
        self.historico.append({...})

        if severidade >= 2:
            # Escalate to higher security level temporarily

            novo_nivel = min(MILITARY, self.tav.nivel_base + severidade - 1)

            self.tav._escalar_nivel(novo_nivel, timeout_tx=1000)
```

Escalation triggers:

- 3 consecutive MAC failures -> Severity 2 -> Escalate 1 level
- Hardware similarity < 0.5 -> Severity 3 -> Escalate 2 levels
- Hardware similarity < 0.7 -> Severity 1 -> Warning only

De-escalation: After timeout (default 1000 transactions), system returns to base level if no further threats detected.

4.11 Dead-Man Switch

Protection against abandoned or compromised instances:

python

```
def verificar_dead_man(self, limite: int = 10000) -> bool:
    """Verifies dead man switch - erases keys if inactive too long."""

    if self.tx_count_global - self.ultima_tx > limite:
        # Erase master entropy - instance becomes unusable

        self.master_entropy = bytes(len(self.master_entropy))

    return False
```

```
return True
```

Use cases:

- Automatic key destruction if device is stolen and unused
- Protection against forensic extraction from inactive devices
- Configurable threshold (default: 10,000 transactions of inactivity)

4.12 Device Identity and Digital Fingerprinting

A unique property of TAV is that each instance's state serves as a **device fingerprint**.

4.12.1 How It Works

Each TAV instance has unique state derived from:

1. **Original seed** (user-controlled secret)
2. **Transaction history** (diverges immediately between devices)
3. **Accumulated physical entropy** (hardware-specific timing jitter)
4. **Clock positions** (deterministic but instance-specific)

Even two devices initialized with the same seed will diverge after the first operation due to different physical entropy.

4.12.2 Identity Applications

Fraud Prevention:

- Verify content authenticity (videos, documents, messages)
- Detect impersonation attempts (state won't match)
- Bind digital assets to specific devices

AI-Generated Content Detection:

- Authentic content carries creator's TAV state signature
- AI-generated content cannot reproduce valid state evolution
- Provides provenance chain for digital media

IoT Device Authentication:

- Pair devices using shared seed
- Verify device identity without network round-trips
- Detect device cloning attempts

4.12.3 Identity vs. Signatures

Use Case	Identity (State)	Signatures
"Is this the same device?"	Yes	Overkill
"Did this device create this content?"	Yes	Yes
"Legally binding transaction"	No	Yes
"Non-repudiation in court"	No	Yes
Quick verification	Yes (fast)	Slower
Computational cost	Low	Higher

Recommendation: Use identity for lightweight verification, signatures for legally binding operations.

4.12.4 Mitigating Tracking Risks

Concern: If state identifies a device, couldn't it be used for unwanted tracking?

Mitigations:

1. **Ephemeral sessions:** Derive temporary states for public interactions
2. **State compartmentalization:** Different instance_ids for different contexts
3. **Voluntary disclosure:** State proof only revealed when user chooses
4. **No central registry:** No server correlating states (unlike IP addresses)

The user controls when and to whom their identity is revealed.

4.12.5 Device Pairing and Synchronization

TAV supports two methods for establishing secure communication between devices:

Method 1: Shared Seed (Pre-Paired)

- Both devices initialized with same seed but different instance_id

- Derive shared secret from common master_entropy
- Each device maintains independent state evolution
- Suitable for: IoT sensor networks, paired personal devices

python

```
# Device A
tav_a = TAV("shared secret phrase", instance_id="device_a")
```

```
# Device B
```

```
tav_b = TAV("shared secret phrase", instance_id="device_b")
```

```
# Both can derive same shared key
```

```
shared_key_a = tav_a._derivar_chave()
shared_key_b = tav_b._derivar_chave()
# shared_key_a == shared_key_b (if same tx_count)
```

Method 2: Key Exchange (Dynamic Pairing)

- One device shares public commitment
- Devices exchange encrypted challenges
- Establish session key without revealing master secrets
- Suitable for: New device pairing, temporary connections

State Synchronization:

- Stateful signatures require verifier to know signer's tx_count
- Options: Include tx_count in signature (implemented), or maintain tx_count database
- Checkpoint can be shared for state recovery (encrypted, requires seed)

4.13 Seed Management and Recovery

4.13.1 User Responsibility

The seed is the root of all TAV security. Like Bitcoin/Ethereum mnemonics:

- **User owns and controls their seed**
- **No central authority can recover it**
- **Loss of seed = loss of identity**

This is a feature, not a bug: no third party can compromise your identity.

4.13.2 Backup Recommendations

Physical Backup (Recommended):

- Write seed on paper
- Store in secure location (safe, safety deposit box)
- Consider multiple copies in different locations

Split Backup (Advanced):

- Use Shamir Secret Sharing to split seed into N parts
- Require K parts to reconstruct (e.g., 3-of-5)
- Distribute parts to trusted parties/locations

What NOT to do:

- Store seed in cloud storage unencrypted
- Take photos of seed with phone
- Email seed to yourself
- Store seed on internet-connected device without additional encryption

4.13.3 Recovery Process

If device is lost but seed is preserved:

1. Initialize new TAV instance with same seed
2. State will restart from transaction 0
3. Old signatures remain valid (verifiable with public key)
4. New instance is a "fresh start" with same identity root

Note: Transaction history is lost. For critical applications, consider periodic checkpoint backups to secure offline storage.

5. Security Analysis

5.1 Threat Model

TAV is designed to resist:

- **Ciphertext-only attacks:** Attacker has access only to encrypted messages
- **Known-plaintext attacks:** Attacker has plaintext-ciphertext pairs
- **Chosen-plaintext attacks:** Attacker can encrypt messages of their choice
- **Replay attacks:** Prevented by unique nonces (timing + monotonic counter)
- **Tampering attacks:** Detected by MAC-Feistel authentication

TAV does **not** claim resistance to:

- **Side-channel attacks:** Timing, power analysis, electromagnetic emanations
- **Fault injection:** Hardware manipulation to induce errors
- **Implementation-specific vulnerabilities:** Bugs in specific codebases

5.2 Entropy Source Security

The physical entropy generator relies on CPU timing jitter, which arises from:

1. **Hardware variability:** Manufacturing variations in transistors
2. **Cache effects:** State-dependent memory access latency
3. **Pipeline effects:** Branch prediction, speculative execution
4. **OS effects:** Scheduler interruptions, context switches
5. **Environmental:** Temperature, voltage fluctuations

Jitterentropy [8] has demonstrated that this source provides sufficient entropy for SP800-90B compliance when properly conditioned. TAV uses similar collection but different conditioning (Feistel mixer vs. SHA3-256).

Measured entropy quality (from 40-test validation suite):

- Shannon entropy: 7.996 bits/byte (theoretical maximum: 8.0)
- Bit bias: 0.15% average (maximum acceptable: 3%)
- Autocorrelation: 0.001 (effective independence)

5.3 Post-Quantum Considerations

TAV's potential quantum resistance derives not from provable hardness assumptions but from the absence of exploitable mathematical structure:

Shor's Algorithm requires algebraic structure (groups with efficiently computable operations) to factor integers or compute discrete logarithms. TAV has no such structureoperations are bitwise logic on physical measurements.

Grover's Algorithm provides quadratic speedup for unstructured search, reducing n-bit security to n/2 bits. TAV with 256-bit keys would have 128-bit security against Grover's search, comparable to AES-256.

Important caveat: This is not a formal security proof. It is an informal argument that quantum algorithms lack obvious attack vectors. A future cryptanalytic breakthrough could change this assessment.

5.4 Comparison with Established Systems

Criterion	Jitterentropy	ChaCha20	ML-KEM	TAV V9.1
Type	Entropy source	Stream cipher	KEM	Complete system
Generates entropy	Yes	No	No	Yes
Stream cipher	No	Yes	No	Yes
MAC included	No	No (Poly1305 separate)	No	Yes
External hash	SHA3-256	None	SHA3/SHAKE	None
Modular addition	N/A	Yes	Yes	No
State model	Entropy pool	Stateless	Stateless	Multi-clock stateful
Standardization	SP800-90B	RFC 8439	FIPS 203	None
Independent audits	Yes	Extensive	Extensive	None
Code size	~2,000 lines	~500 lines	~5,000+ lines	~1,600 lines (C)
Performance	Entropy only	~1,000 MB/s	<1ms keygen	~3.9 MB/s
Key size	N/A	256 bits	800-1568 bytes	16-32 bytes

Signature size	N/A	N/A	2,420-4,627 bytes	66-72 bytes
-----------------------	-----	-----	-------------------	-------------

5.5 Known Limitations

We explicitly acknowledge the following limitations:

1. **No formal security proofs:** TAV's security is not reducible to established hard problems. We cannot prove resistance to unknown attacks.
2. **No independent audit:** The code has not been reviewed by external cryptographers. Implementation bugs may exist.
3. **No side-channel analysis:** Timing variations in TAV operations have not been analyzed. The implementation is not guaranteed constant-time.
4. **Limited cryptanalysis:** TAV has not been subjected to extensive attack by the cryptographic community.
5. **Application scope:** TAV is designed for scenarios where simplicity, auditability, and self-contained operation matter more than maximum throughput. It is ideal for IoT, low-bandwidth communications, and key storage not for high-volume encryption like video streaming or disk backup.
6. **Timer characteristics:** Entropy quality benefits from high-resolution timers when available. On platforms with low-resolution timers, TAV compensates through multiple measurements and state accumulation. Note that TAV **evolves state on demand**, not continuously a sleeping device consumes no entropy resources.

6. Implementation

7.1 Reference Implementations

Complete implementations are provided in four languages:

Language	Lines of Code	Features	Primary Use Case
Python	1,551	Complete reference implementation	Prototyping, testing
C	1,608	Full system + OpenSSL engine	Embedded, systems
Rust	1,101	Memory-safe implementation	Security-critical
JavaScript	928	Browser and Node.js support	Web applications
Arduino	790	ESP32/ESP8266 optimized	IoT devices

Performance Optimizations

All implementations include:

1. **Lookup tables for rotation:** 8 tables of 256 bytes precomputed
2. **Primes as bytes:** 1,464 primes already converted to binary format
3. **Entropy pool with TTL:** Entropy discarded after 1,000 transactions (transaction-based, not wall-clock time)
4. **Tiered threat verification:** IoT=20 tx, Consumer=10 tx, Enterprise=5 tx, Military=1 tx

All implementations:

- Produce identical outputs for identical inputs (validated by test vectors)
- Support all four security levels
- Are self-contained with no external dependencies

7.2 Cross-Platform Interoperability

We validated interoperability through bidirectional test vector exchange:

Test procedure:

1. Generate test vectors in C (hash, encrypt operations)
2. Export vectors as JSON
3. Import and verify in JavaScript
4. Generate new vectors in JavaScript
5. Import and verify in C

Results:

- Hash compatibility: 100% (1,000 vectors tested)
- Encrypt/decrypt round-trip: 100% (10,000 operations)

7.3 Code Quality

The reference implementations follow secure coding practices:

- **No dynamic memory allocation** in critical paths (C implementation)
- **Bounds checking** on all array accesses

- **Input validation** on all public API functions
- **Comprehensive test coverage** (40 test categories)
- **Documentation** for all public interfaces

7. Performance Evaluation

7.1 Test Environment

All tests were conducted on:

- **Platform:** Linux (Ubuntu 22.04)
- **CPU:** Intel Xeon (representative of server deployments)
- **Language:** Python reference implementation (for consistency across tests)

7.2 The Throughput Comparison Problem

Critical Note: Why direct comparisons are misleading

When reading "ChaCha20 does 1,000 MB/s" and "TAV does 600 KB/s", it appears TAV is $\sim 1,600$ slower. **This is fundamentally incorrect** because the metrics measure different operations:

What "1,000 MB/s" for ChaCha20 measures	What "600 KB/s" for TAV measures
Receives already derived key	Collects physical entropy from CPU
Receives already generated nonce	Evolves state of 4 clocks
Encrypts data	Derives key from current state
Computes MAC	Generates guaranteed unique nonce
2 operations	Encrypts data

7.3 Fair Comparison: Isolating Components

To compare properly, we measured each TAV component in isolation:

Time Decomposition per Operation (encrypt 1KB, Consumer level)

Component	Time	% of Total	Comparable to
Derive key	0.007 ms	0.4%	HKDF
Generate nonce	0.004 ms	0.2%	CSPRNG

Generate keystream	0.140 ms	8.8%	ChaCha20
XOR encrypt	0.045 ms	2.8%	Stream cipher
Compute MAC	1.403 ms	87.8%	Poly1305

Key finding: TAV's keystream, in isolation, achieves ~5 MB/s in Python comparable to ChaCha20 in pure Python implementation (~3-5 MB/s).

Throughput by Component (1MB messages)

Component	TAV (Python)	ChaCha20 (pure Python)	ChaCha20 (C/AVX2)
Keystream	~5.2 MB/s	~3-5 MB/s	~1,000 MB/s
Complete system	~0.6 MB/s	N/A	N/A

7.4 Complete System Throughput

By Security Level (1KB messages)

Level	Throughput	Tick Latency	Verification every
IoT	~850 KB/s	~2 s	20 transactions
Consumer	~620 KB/s	~4 s	10 transactions
Enterprise	~480 KB/s	~5 s	5 transactions
Military	~490 KB/s	~4 s	1 transaction

By Message Size (Consumer level)

Size	Throughput	Observation
64 B	~320 KB/s	Fixed overhead dominates
256 B	~520 KB/s	Typical balance
1 KB	~620 KB/s	Common case
4 KB	~640 KB/s	Near maximum
16 KB+	~650 KB/s	MAC-limited

7.5 Bottleneck Analysis: MAC-Feistel

MAC-Feistel consumes ~88% of total time because:

1. Processes all data in 32-byte blocks
2. Applies multiple Feistel rounds per block
3. Uses logic-only operations (no specialized instructions)

Design trade-off: Security and auditability vs. raw throughput.

7.6 Operations Per Byte Analysis

Elementary operation count:

Algorithm	Ops/byte	Operation types
TAV Keystream	~6	XOR, AND, OR, ROT (lookup)
TAV MAC-Feistel	~1.3 (amortized)	XOR, AND, OR, ROT
TAV Total	~7.3	Logic-only
ChaCha20	~20	ADD, XOR, ROT
Poly1305	~3-6	MUL 128-bit mod p
ChaCha20-Poly1305	~23-26	ADD, XOR, ROT, MUL

Key finding: TAV uses **3-4 fewer operations** per byte than ChaCha20-Poly1305!

Why is measured ChaCha20 faster?

1. **SIMD:** ChaCha20 with AVX2 processes 4-8 blocks in parallel
2. **Counter mode:** Each block is independent (perfect parallelization)
3. **Decades of optimization:** Highly refined implementations

Estimated throughput in C:

Implementation	Estimated Throughput
TAV Python	~0.6 MB/s (interpreter overhead)
TAV C basic	~100-300 MB/s
TAV C with SIMD	~300-800 MB/s
ChaCha20-Poly1305 (AVX2)	~1-3 GB/s

Real performance ratio: ~3-10 (not ~1000 as appears in Python)

TAV is intrinsically efficient because:

- No multiplication (Poly1305 uses it, ~50-100 cycles)

- No cryptographic hash (HKDF uses SHA256)
- All operations are ~1 CPU cycle
- RDTSC (perf_counter_ns) is just 1 instruction

7.7 Operating Model: On-Demand

TAV **does not** consume resources when the device is idle:

Typical IoT device:

[Sleeping 99.9% of time] TAV cost: ZERO

[Wakes to transmit] Collects entropy + encrypts + authenticates

[Returns to sleep] TAV cost: ZERO

This contrasts with systems that:

- Maintain active entropy pools
- Require periodic reseeding
- Consume CPU cycles in background

7.8 Practical Capacity

What ~600 KB/s means in practice:

Scenario	Messages/sec	Sufficient?
IoT sensor (100 bytes)	6,000	Yes - Far beyond
Automation command	1,000+	Yes - Adequate
Encrypted chat (1KB)	600	Yes - Adequate
Video streaming	<1 frame	No - Inadequate
Disk backup	<1%	No - Inadequate

7.9 When TAV Is Appropriate

Yes Ideal use cases:

- Devices without reliable hardware RNG

- Auditability requirements (simple code, no dependencies)
- Low-bandwidth applications (< 1,000 msgs/s)
- Memory-constrained devices
- Environments where independence from external libraries is critical

No Inappropriate use cases:

- High-volume encryption (video, backups)
- FIPS/NIST compliance requirements
- Environments with available hardware acceleration (AES-NI)
- Formal security proofs required

7.10 Comparison with NIST PQC

For reference, ML-KEM and ML-DSA data (FIPS 203/204):

ML-KEM (Kyber) - Key Exchange on x86_64 with AVX2:

Operation	ML-KEM-768	ECC P-384	RSA-7680
Key Generation	~3.2 s	~9.3 s	~66 seconds
Encapsulation	~0.09 s	~6.6 s	~2.3 s
Decapsulation	~0.1 s	~6.6 s	~332 ms

Observation: ML-KEM is extremely efficient. The bottleneck for PQC adoption is not key exchange--it's signatures (ML-DSA is 10-15x slower than ECDSA).

7.11 Why NIST PQC Is Not Widely Used Yet

Despite ML-KEM's excellent performance, adoption remains slow:

1. Very Recent Standards

- FIPS 203/204/205 finalized: August 2024 (only ~1 year ago!)
- Before that, algorithms were "candidates"--no one deploys candidates in production

2. Cryptographic Libraries Still Catching Up

Library	ML-KEM/ML-DSA Status
OpenSSL	Support in 3.5 (April 2025) - very new
libsodium	"On roadmap" - no timeline

MbedTLS (IoT)	Planned - not implemented
LibreSSL	No announced plans
BoringSSL (Google)	Yes Supports (only one in real production)
wolfSSL	Yes Supports

3. Legacy Infrastructure

- 75% of SSH servers run versions from 2015-2022
- Less than 20% of TLS servers use TLS 1.3 (only version supporting PQC)

4. Compatibility Problems Already Occurred

- Chrome 124 enabled hybrid Kyber
- Rollback in Chrome 131 due to "widespread TLS handshake failures"
- Corporate firewalls blocking larger handshakes

5. Much Larger Sizes

Metric	ECC (current)	ML-KEM/ML-DSA
Public key	32-64 bytes	1,184-1,952 bytes
Signature	64-72 bytes	2,420-4,627 bytes
TLS overhead	~100 bytes	~2,000+ bytes

6. Cryptographic Transitions Are Slow

- SHA-1 -> SHA-2: vulnerabilities in 2005, full deprecation 2017+
- More than 12 years for a "simple" transition
- NIST recommends PQC migration by 2035

7.12 The Stateful Advantage for IoT

For a constrained IoT device to perform authenticated post-quantum encryption:

PQC Stack (stateless): ML-KEM key exchange (requires network round-trip: 2,272 bytes transmitted + 50-200ms latency) -> SHAKE256 KDF -> ChaCha20 cipher -> Poly1305 MAC -> External entropy source. Total: 5+ libraries, multiple dependencies.

TAV Stack (stateful): TAV.encrypt() -> Done. Total: 1 system, zero dependencies, zero network communication for key establishment.

7.13 Key Derivation: Fair Comparison

Isolating the 'obtain a key ready for use' operation:

Operation	ML-KEM-768	TAV
CPU time for key	~3.4 us	~7 us
Bytes transmitted	2,272 bytes	0 bytes
Network latency (LAN)	~1 ms	0
Network latency (Internet)	~50-200 ms	0
Total real-world time	~50-200 ms	~7 us

Key insight: For IoT scenarios with network communication, TAV is 7,000-28,000x faster in real-world key establishment because it eliminates the network round-trip entirely.

7.14 Resource Comparison for Constrained Devices

Resource	PQC Stack	TAV
Code size	~8,500+ lines	~1,600 lines (C)
RAM usage	~10+ KB	~500 bytes
Public key size	1,184-1,952 bytes	16-32 bytes
Signature size	2,420-4,627 bytes	N/A (future work)
External dependencies	SHA3, SHAKE, RNG	None
Fits in 64KB flash?	Difficult	Yes
Runs in 32KB RAM?	No	Yes

7.15 TAV's Niche

TAV provides an alternative approach for specific scenarios:

- Systems without SHA3:** Devices that cannot implement SHA3/SHAKE256
- Full auditability:** 100% auditable code with no external dependencies

3. **Self-contained operation:** No need to combine multiple libraries
4. **Integrated entropy:** Built-in physical entropy collection
5. **Education/Research:** Demonstration of alternative cryptographic concepts

TAV is appropriate for:

- Low-bandwidth IoT sensors (< 100 messages/second)
- Applications requiring code auditability
- Research and educational purposes
- Devices without hardware RNG

TAV is NOT appropriate for:

- Applications requiring FIPS certification
- Systems requiring formal security proofs
- Regulated critical infrastructure
- High-volume bulk encryption (use ChaCha20/AES)
- Digital signatures (use ML-DSA or ECDSA)

8. Validation Results

8.1 Test Suite Overview

TAV V9.1 was subjected to a comprehensive validation suite of **40 tests** across **14 categories**:

Category	Tests	Result
Statistical (NIST SP 800-22)	3	3/3 PASS
Entropy (SP 800-90B alignment)	3	3/3 PASS
Advanced statistical	2	2/2 PASS
TAV-specific	3	3/3 PASS
MAC-Feistel	2	2/2 PASS
Nonce uniqueness	1	1/1 PASS
Performance	2	2/2 PASS
Stress testing	4	4/4 PASS

Long-duration	3	3/3 PASS
Edge cases	6	6/6 PASS
Prime/box verification	3	3/3 PASS
Mixer analysis	2	2/2 PASS
Pattern detection	3	3/3 PASS
Security properties	3	3/3 PASS
TOTAL	40	40/40 PASS (100%)

8.2 Statistical Test Results (NIST SP 800-22)

Test	Statistic	Criterion	Result
Frequency (Monobit)	p = 0.536	p > 0.01	Yes - PASS
Runs	p = 0.630	p > 0.01	Yes - PASS
Serial (Pairs)	/df = 1.70	< 3.5	Yes - PASS

8.3 Entropy Quality Results

Metric	Value	Criterion
Shannon entropy	7.996 bits/byte	7.9
Compression ratio	1.0004	> 0.95
Bit bias (average)	0.15%	< 3%
Autocorrelation	0.001	< 0.05

Bit-level bias analysis:

Bit 0: 0.18% Bit 4: 0.12%

Bit 1: 0.15% Bit 5: 0.21%

Bit 2: 0.09% Bit 6: 0.14%

Bit 3: 0.19% Bit 7: 0.11%

Maximum: 0.32% Average: 0.15%

8.4 Avalanche Properties

Property	Value	Criterion
Key avalanche	48.26%	35-65% (ideal: 50%)
MAC avalanche	50.78%	35-65% (ideal: 50%)
Mixer diffusion	7.81 bits/byte	7.5

8.5 Stress Test Results

Test	Operations	Errors	Result
Continuous encrypt/decrypt	1,000	0	Yes - PASS
Variable message sizes (0 to 100KB)	8 sizes	0	Yes - PASS
All security levels	4 50 = 200	0	Yes - PASS
Rapid sequential operations	500	0	Yes - PASS

8.6 Security Properties

Property	Test	Result
Nonce uniqueness	5,000 nonces	100% unique
Tamper detection	20 bit-flip positions	100% detected
Seed independence	10 different seeds	100% unique outputs
Replay prevention	100 duplicate attempts	100% rejected

9. Discussion

9.1 Appropriate Use Cases

TAV is suitable for:

1. **Research and experimentation:** Exploring alternative cryptographic paradigms
2. **Education:** Teaching cryptographic concepts with auditable code
3. **IoT prototyping:** Lightweight cryptography for constrained devices
4. **Low-bandwidth communication:** Where 3.9 MB/s is sufficient
5. **Auditability-critical applications:** Where code simplicity enables thorough review

9.2 Inappropriate Use Cases

TAV should **NOT** be used for:

1. **Production systems:** Until independent security analysis is completed
2. **Sensitive data protection:** Where proven security is required
3. **Compliance requirements:** TAV is not FIPS-certified
4. **High-throughput applications:** Where ChaCha20/AES speeds are necessary
5. **Side-channel-sensitive environments:** Without additional hardening

9.3 Future Work

The following areas warrant further investigation:

1. **Formal security analysis:** Prove or disprove security properties
2. **Cryptanalytic attack:** Attempt to break TAV under various attack models
3. **Side-channel hardening:** Implement constant-time operations
4. **Performance optimization:** SIMD, hardware acceleration
5. **Extended test suites:** Full NIST SP 800-22 (all 15 tests), Dieharder, TestU01
6. **Protocol integration:** TLS integration, key exchange protocols

9.4 Call for Community Analysis

We invite the cryptographic community to:

- **Analyze:** Examine the design for weaknesses
- **Attack:** Attempt to break the security claims
- **Critique:** Identify flaws in reasoning or implementation
- **Improve:** Suggest enhancements while maintaining the core philosophy
- **Verify:** Independently reproduce our test results

The complete source code, test vectors, and documentation are available at:

<https://github.com/carlostbastos/tav-crypto>

10. Conclusion

We have presented TAV (Transactional Asynchronous Verification) version 0.9, a comprehensive cryptographic system that integrates physical entropy collection, stateful stream cipher, message authentication, digital signatures, and device identity into a unified architecture using only logic operations.

Key contributions:

1. **The concept of 'ephemeral structure'**--algebraic elements that exist but change too rapidly to be exploited, providing a different security model than both static structure and pure chaos.
2. **Novel multi-clock transactional architecture** with mathematically guaranteed state cycle lengths of 563,327 transactions minimum.
3. **Complete elimination of modular addition**, avoiding rotational cryptanalysis vulnerabilities.
4. **Compact digital signatures** (66-72 bytes) that are 36-70x smaller than post-quantum ML-DSA, with certificate support for identity management.
5. **Device identity framework** enabling fraud prevention, content authenticity verification, and detection of AI-generated content through unique state fingerprints.
6. **Threat management system** with automatic security escalation and dead-man switch protection.
7. **Automatic encrypted checkpointing** for state persistence and recovery with hardware change detection.
8. **Demonstrated advantages for constrained IoT devices:** 7,000-28,000x faster key establishment than PQC in networked scenarios, ~500 bytes RAM vs ~10KB+, and complete self-containment.
9. **Comprehensive validation** across 40 tests with 100% pass rate.
10. **Cross-platform implementations** in five languages with verified interoperability.

Key limitations:

1. No formal security proofs
2. No independent audit
3. Slower than SIMD-optimized conventional ciphers for bulk encryption
4. No side-channel analysis
5. Stateful signatures require careful key management

TAV represents an exploration of physics-pervasive cryptography with ephemeral structure--where physical phenomena govern the entire cryptographic operation, and algebraic elements exist but cannot be exploited due to continuous evolution. The addition of digital signatures and device identity capabilities extends TAV from a cipher into a complete identity and authentication framework.

Whether this approach can achieve security comparable to mathematically-grounded constructions remains an open question for the community. We believe TAV's unique properties--particularly device fingerprinting and compact signatures--offer valuable tools for fraud prevention and content authenticity in an era of sophisticated AI-generated content.

We welcome scrutiny, criticism, and collaboration.

Acknowledgments

This work was developed independently without institutional affiliation or external funding. The author thanks the open-source community for tools and infrastructure that made this research possible.

References

- [1] NIST, "Post-Quantum Cryptography Standardization," FIPS 203, 204, 205, August 2024.
Available: <https://csrc.nist.gov/projects/post-quantum-cryptography>
- [2] L. Chen et al., "NIST Post-Quantum Cryptography Standardization Process: Third Round Candidate Announcement," NISTIR 8413, July 2022.
- [3] D. Khovratovich and I. Nikolic, "Rotational Cryptanalysis of ARX," Fast Software Encryption (FSE), 2010.
- [4] D. J. Bernstein, "ChaCha, a variant of Salsa20," Workshop Record of SASC, 2008.
- [5] D. J. Bernstein, "The Salsa20 family of stream ciphers," New Stream Cipher Designs, Springer, 2008.
- [6] J.-P. Aumasson et al., "SHA-3 proposal BLAKE," Submission to NIST, 2008.
- [7] N. Ferguson et al., "The Skein Hash Function Family," Submission to NIST, 2008.
- [8] S. Mller, "CPU Time Jitter Based Non-Physical True Random Number Generator," 2014.
Available: <https://www.chronox.de/jent.html>
- [9] NIST, "Recommendation for the Entropy Sources Used for Random Bit Generation," SP 800-90B, January 2018.
- [10] T. Ts'o, "random: use a LFSR for jitter entropy generation," Linux kernel commit, 2019.

- [11] NIST, "NIST IR 8547: Transition to Post-Quantum Cryptography Standards, November 2024.
- [12] M. Bellare, T. Kohno, and V. Shoup, "Stateful Public-Key Cryptosystems: How to Encrypt with One 160-bit Exponentiation," ACM CCS, 2006.
- [13] P. W. Shor, "Algorithms for quantum computation: discrete logarithms and factoring," IEEE FOCS, 1994.
- [14] L. K. Grover, "A fast quantum mechanical algorithm for database search," STOC, 1996.

Appendix A: Test Vectors

A.1 Hash Test Vector

```
json
{
  "input": "TAV Test Vector",
  "input_hex": "544156205465737420566563746f72",
  "level": "CONSUMER",
  "expected_hash_hex": "a7c4e2f8b3d1...[32 bytes]"
}
```

A.2 Encryption Test Vector

```
json
{
  "plaintext": "Encrypt this message",
  "plaintext_hex": "456e63727970742074686973206d657373616765",
  "level": "CONSUMER",
  "ciphertext_length": "60 bytes (20 + 40 overhead)",
  "mac_verified": true
}
```

Complete test vectors available at: <https://github.com/carlostbastos/tav-crypto/tests/>

Appendix B: Security Level Configuration

```
python
CONFIG_BY_LEVEL = {
    "IoT": {
        "master_entropy": 32,
        "key_bits": 128,
        "key_bytes": 16,
        "mac_bytes": 8,
        "nonce_bytes": 8,
        "n_xor": 2,
        "mixer_rounds": 2,
        "mac_rounds": 4,
        "initial_boxes": [1, 2],
        "available_boxes": [3],
    },
    "Consumer": {
        "master_entropy": 48,
        "key_bits": 192,
        "key_bytes": 24,
        "mac_bytes": 12,
        "nonce_bytes": 12,
        "n_xor": 2,
        "mixer_rounds": 3,
        "mac_rounds": 6,
        "initial_boxes": [1, 2, 3],
        "available_boxes": [4],
    },
    "Enterprise": {
        "master_entropy": 64,
```

```
"key_bits": 256,  
"key_bytes": 32,  
"mac_bytes": 16,  
"nonce_bytes": 16,  
"n_xor": 3,  
"mixer_rounds": 4,  
"mac_rounds": 8,  
"initial_boxes": [1, 2, 3, 4],  
"available_boxes": [5],  
,  
"Military": {  
"master_entropy": 64,  
"key_bits": 256,  
"key_bytes": 32,  
"mac_bytes": 16,  
"nonce_bytes": 16,  
"n_xor": 4,  
"mixer_rounds": 6,  
"mac_rounds": 8,  
"initial_boxes": [1, 2, 3, 4, 5],  
"available_boxes": [6],  
,  
}  
}
```

Appendix C: Licensing

TAV is released under a dual license:

Open Source: GNU Affero General Public License v3.0 (AGPL-3.0)

Commercial: Free commercial use until May 2027. After this date, commercial users must either comply with AGPL-3.0 or obtain a commercial license.

See COMMERCIAL_LICENSE.md in the repository for details.

How to Cite

bibtex

```
@software{bastos2025tav,  
author = {Bastos, Carlos Alberto Terencio de},  
title = {{TAV Clock Cryptography: Stateful Encryption with  
Ephemeral Structure and Device Identity}},  
version = {0.9},  
year = {2025},  
month = {November},  
url = {https://github.com/carlostbastos/tav-crypto},  
note = {Experimental research software - Initial public release}  
}
```

TAV Clock Cryptography v0.9

*Transactional Asynchronous Verification
November 2025 - Initial Public Release*